

A Systematic Trading Approach from Data Mining to Live Deployment

Documentation

Submitted by:

Juri Stoffers

Supervised by:

Dr. Geoffrey Ostrin

July 19, 2025

Contents

1	Orderbook Delta defenition	4
2	Outlier Detection Using Mean and Standard Deviation (Z-Score Based Outlier Detection)	5
2.1	Normal Range	5
2.2	My Hypothesis	5
2.3	Normal Range	5
2.4	Outlier Condition	6
2.5	Idea behind	7
2.6	Future Plans	7
3	Measuring Volatility After Price Outlier Detection	8
3.1	Dictionary of Terms	8
3.2	In Python code	8
3.3	Statistical evidence	9
4	Optimising for best Z-Score thresold for outliers	10
5	Measuring avearge return after price outlier detection	11
5.1	Formulars	11
5.2	Dictionary of Terms	11
6	Underlying strategy Bias	12
7	Finding an edge	14
7.1	Fees	14
7.2	No Fees results	14
7.3	With Fees results	15
7.4	No clear path	16
8	Comparison of Delta Indicators on 39,694-minute sample	17
9	Combining Indicators	18
10	Mapping Orderbook delta EMA with standard deviation shifing	19
11	Formulars	19
12	From outliers to strategy	20
12.1	Progress description	20
12.2	Where did I start my research	20
12.3	Finding an edge	21

12.4 Monte Carlo Testing	22
12.5 Backtesting framework	24
13 Developing a strategy	26
13.1 Initial thoughts	26
13.2 Reverse strategy	26
13.3 First strategy test	26
13.3.1 Code	26
13.3.2 Strategy 1 results	27
13.4 Explaining figure	28
13.5 What did change	28
14 sources	29
14.1 How good or random is your trading	29
14.2 Testing 1000 strategies at the same time (via Monte Carlo)	29
14.3 Tweets	29
14.4 TRDR This platform allow you to use different kind of metrics on different Timeseries datasets (BTC/USD Price, orderbookdelta and Open Interest	29
14.5 Trend line automation	29
14.6 Vectorbt github documentation respository.	29
14.7 Backtesting	29

1 Orderbook Delta defenition

The orderbook delta is calculate by the difference between the sum of the bid and ask orders at a certain depth. Formular:

$$\Delta_{x\%} = \sum_{i=1}^{x\%} \Delta_i$$

- $\Delta_{x\%}$ is the sum of the orderbook delta for the last $x\%$ of the orderbook.
- Δ_i is the orderbook delta for the i -th level of the orderbook.

2 Outlier Detection Using Mean and Standard Deviation (Z-Score Based Outlier Detection)

2.1 Normal Range

What I want to test is how price reacts to anomalous orderbook delta movements, particularly in scenarios where unrealistic or clearly outlying values are detected. In cryptocurrency markets, such inefficiencies can be caused by various events, one example is liquidation events that interact with passive demand order stacked zones. During these events, the orderbook delta exhibits significant increases, providing a clear signal of market stress. This research will focus on understanding the relationship between rapid delta movements and how price reacts after these events.

2.2 My Hypothesis

- I expect realized volatility to increase after an outlier is detected.
- I expect a return to the mean after a strong outlier is detected.

2.3 Normal Range

$$\mu(\Delta) \pm 2\sigma(\Delta)$$

This means most data points (about 95% if normally distributed) are expected to lie within this range.

2.4 Outlier Condition

A value is considered an outlier if:

$$\Delta < \mu(\Delta) - 2\sigma(\Delta) \quad \text{or} \quad \Delta > \mu(\Delta) + 2\sigma(\Delta) \quad (1)$$

This is a simple Z-score based outlier detection.

- Δ - Orderbook Delta Depth with a certain depth I will test on: $\Delta_{1\%}$ $\Delta_{2.5\%}$ $\Delta_{5\%}$ from Coinbase (BTC/USD)
- This basically means we take a delta of the Bid and Ask orders which are in a range of x% from the current price.
- $\mu(\Delta)$ — Mean of the last 1440 values of Δ before time t
- $\sigma(\Delta)$ - Standard deviation over the last 1440 Δ values before time t

2.5 Idea behind

- This method assumes data is roughly normally distributed.
- Using 2σ captures approximately 95% of data points under a normal distribution.
- You can adjust the multiplier (e.g., 3σ) for stricter or looser thresholds.

2.6 Future Plans

- Test on more data
- use rolling windows (e.g. 1 day or 1 week) for local context.
- Compare sensitivity with $\pm 1.5\sigma$ or $\pm 2.5\sigma \rightarrow$ optimize for best results

3 Measuring Volatility After Price Outlier Detection

$$r_t = \frac{P_t - P_{t-1}}{P_{t-1}} \quad (2)$$

3.1 Dictionary of Terms

- P_t Asset price at time t .
- $r_t = \frac{P_t - P_{t-1}}{P_{t-1}}$ – 1-minute price return at time t .
- $\sigma_t^{(15)}$ – Realized volatility: the standard deviation of the next 15 one-minute returns,

$$\sigma_t^{(15)} = \sqrt{\frac{1}{14} \sum_{i=1}^{15} (r_{t+i} - \bar{r}_t)^2}, \quad \bar{r}_t = \frac{1}{15} \sum_{i=1}^{15} r_{t+i}. \quad (3)$$

aligned so that at time t it measures volatility over $t + 1$ to $t + 15$.

3.2 In Python code

```
import pandas as pd
df = pd.read_csv(file_path)
df.set_index('timestamp', inplace=True)
#Compute 1-min return of delta_5

df['r_t'] = df['price'].pct_change().fillna(0)

#compute rolling std of the future 15 min window

window = 15

#rolling on r_t, then shift forward so index t hold vol of t+1...t+15
df['future_vol_15'] = (
    df['r_t']
    .rolling(window=window)
    .std()
    .shift(-window)
)
```

3.3 Statistical evidence

Once an outlier is detected (1) inside of the Orderbook Δ , we calculate the 15-minute ahead realized volatility using Equation: (3)

if a Δ_t values is flagged as an outlier (1) we record

$$\sigma_t^{(15)} = \sqrt{\frac{1}{14} \sum_{i=1}^{15} (r_{t+i} - \bar{r}_t)^2},$$

We then form two samples over our full dataset which during this test includes 104 957 one minutes intervals of P and Orderbook Δ :

$$\mathcal{S}_{\text{out}} = \{\sigma_t^{(15)} : t \text{ is an outlier}\}, \quad \mathcal{S}_{\text{non}} = \{\sigma_t^{(15)} : t \text{ is not an outlier}\}.$$

Sample mean results:

$$\bar{\sigma}_{\text{out}}^{(15)} = 0.0006244, \quad \bar{\sigma}_{\text{non}}^{(15)} = 0.0005138,$$

This concludes an increase of r_t of roughly 21.5%

To check Statistical evidence

- a two-sample *t*-test (unequal variances), which yields

$$T = 24.72, \quad p = 4.79 \times 10^{-132},$$

- a Mann–Whitney *U*-test, which returns

$$p = 4.02 \times 10^{-157}.$$

4 Optimising for best Z-Score threshold for outliers

As state inside of (1) we use a Z-Score threshold of 2 to detect outliers. I now want to see if by any chance there is a better Threshold value

To compare the outliers Volatility with the non outliers volatility I will use the following formular:

$$U(z) = \frac{\overline{\sigma}_{\text{out}}^{(15)}}{\overline{\sigma}_{\text{non}}^{(15)}} \quad (4)$$

First I run an optimization for the thresholds of the Z-Score to find the best threshold value for the outliers on a 45 days dataset. After that I compare the result with a 107880 minutes dataset. Where I also run an optimization for the thresholds of the Z-Score to find the best threshold value for the outliers.

Top 3 z -values with largest volatility uplift 69811-minutes sample

z	N_{out}	$U(z)$ (%)	Mann-Whitney p
3.8	221	+58.36	1.913×10^{-51}
3.9	166	+61.99	4.227×10^{-41}
4.0	117	+58.36	8.228×10^{-28}

Same z -values on extended dataset (Walk forward $v1$)

z	N_{out}	$U(z)$ (%)	Mann-Whitney p
3.8	644	+51.73	2.549×10^{-77}
3.9	561	+53.28	6.150×10^{-88}
4.0	479	+50.90	5.517×10^{-59}

Top 3 z -values with largest volatility uplift (Walk forward $v2$)

z	N_{out}	$U(z)$ (%)	Mann-Whitney p
4.8	214	+65.10	6.475×10^{-33}
4.9	197	+66.81	3.693×10^{-29}
5.0	181	+70.53	6.599×10^{-28}

5 Measuring average return after price outlier detection

5.1 Formulas

Once a Δ_t Outlier is detected we calculate the 15-min forward return of BTC/USD price

$$\text{Ret}_t^{(15)} = \frac{P_{t+15} - P_t}{P_t} \quad (4)$$

We then differentiate between a bullish and a bearish outlier. Which is already defined (1)

$$\overline{\text{Ret}}_{\text{bull}}^{(15)} = \frac{1}{|\mathcal{T}_{\text{bull}}|} \sum_{t \in \mathcal{T}_{\text{bull}}} \text{Ret}_t^{(15)} \quad (7)$$

$$\overline{\text{Ret}}_{\text{bear}}^{(15)} = \frac{1}{|\mathcal{T}_{\text{bear}}|} \sum_{t \in \mathcal{T}_{\text{bear}}} \text{Ret}_t^{(15)} \quad (8)$$

5.2 Dictionary of Terms

- Price at a certain time: P_t
- 15-min forward return: $\text{Ret}_t^{(15)}$

#compute 15-min forward return of BTC/USD price

6 Underlying strategy Bias

Every single parameter has to fight to be implemented into my strategy. To get some kind of filter since we are working with an asset which has clear trends so it isn't stationary we need to do some trend identification. Different trends are also called different regimes. I'll call it the underlying bias. Some simple examples are for a bias are:

- $\Delta_{5\%} < 0$ (more passive demand than supply)
- $EMA_n > EMA_x$ (Crossing of two EMAs with different time periods n and x)

But I want to get some clear trend identification where we divide into the three different categories:

- Uptrend
- Downtrend
- Ranging

Uptrend is defined as a period where the price is making higher highs and higher lows. Downtrend is defined as a period where the price is making lower highs and lower lows. Ranging is defined as a period where price is not making higher highs and higher lows or lower highs and lower lows. But for my approach I will just use this if we are getting mixed signals.

Ideas I want to test on:

- Market Structure definition (Higher Highs and Higher Lows)
- Relative strength index on Price and $\Delta_x\%$

7 Finding an edge

7.1 Fees

So the first problem you run into if building a strategy is the fees. You have to pay fees for every trade you make. A fee is a fixed percentage amount of the trade size you take. On the exchange I choose, Hyperliquid the fee is about 0.003% of the trade size. Which is a normal amount inside the crypto space. Just to showcase how important fees are I will show you a simple example.

Inside of my Market simulation where I can test different kind of strategies I will set up a simple strategy. When ever a bullish outlier in the delta appears and the trend function sais that we are inside of an uptrend I will buy Bitcoin and wait for 240 Minutes.

```
#Define a mask for the bullish outliers inside of an uptrend
outlier_mask = (df_temp['outlier_context'] == 's') & (df_temp['trend'] == 'Uptrend')

#Define the entry and exit signals with the Vector BT library
pf = vbt.Portfolio.from_signals(
    close=price,
    entries=entry_signals,
    exits=exit_signals,
    init_cash=100, #Begining amount if cash
    freq='1T'
)
```

7.2 No Fees results

Table 1: No Fees	
Metric	Value
Holding Period	240 minutes
Total Return	2.599%
Mean sharp ratio	3.067088
Total Signals	187

7.3 With Fees results

Table 2: With Fees	
Metric	Value
Holding Period	240 minutes
Total Return	1.537%
Mean sharp	1.841555
Mean Z-Score	2.48

7.4 No clear path

Finding an edge is a very hard task. There is no clear path to success. You kinda have to try by trail and error every error could be a step further but also a potential path into a dead end.

There are three main steps to find an edge.

- Copying other strategies
- Finding an edge by yourself
- Build a systematic version of your discretionary trading.

8 Comparison of Delta Indicators on 39,694-minute sample

Period	$N_{signals}$	$U(z)$ (%)	T-statistic p
Delta2.5 Strategy ($N_{total} = 9,182$)			
5m	9,182	+52.03	1.851×10^{-41}
60m	9,182	+55.94	3.537×10^{-4}
360m	9,182	+59.20	12.524×10^0
Delta5 Strategy ($N_{total} = 12,234$)			
5m	12,234	+51.38	3.382×10^{-7}
60m	12,234	+53.44	4.010×10^{-1}
360m	12,234	+57.04	13.617×10^0

Table 3: Comparison of Delta2.5 vs Delta5 strategies across different time periods

- $N_{signals}$ represents number of trading signals
- $U(z)$ represents positive return ratio
- T-statistic p represents statistical significance

9 Combining Indicators

Here I visualised the swing points, the EMA spread and the 100 outliers with the highest Z-Score in the same plot.




Figure 1: combined indicators png

1

¹Chart made with Matplotlib and Seaborn

10 Mapping Orderbook delta EMA with standard deviation shifing

11 Formulars

$$EMA_t = \alpha \cdot \Delta_t + (1 - \alpha) \cdot EMA_{t-1} \quad (5)$$

Dictionary of Terms

- EMA_t - Exponential Moving Average at time t
- Δ_t - Orderbook delta at time t
- α - Smoothing factor
- $\alpha = \frac{2}{n+1}$ where n is the period of the EMA
- $\sigma(\Delta)$ - Standard deviation of the orderbook delta with a

We'll now shift the EMA by the standard deviation of the orderbook delta

12 From outliers to strategy

12.1 Progress description

I had a pretty hard time going from developing a logic for the outliers of the Delta. Finding that volatility increases after outliers was a good first find which didn't take long to find. But finding some definition condition to meet directional bias was very difficult. By directional bias I mean a price direction which is not random after some kind of condition. It was pretty clear from the beginning that the outliers by themselves won't offer any kind of directional bias.

An outlier is defined as:

- (1). $\Delta < \mu(\Delta) - 2\sigma(\Delta)$ or $\Delta > \mu(\Delta) + 2\sigma(\Delta)$

Then since Bitcoin clearly is not a stationary asset I needed to find some kind of trend identification. I found that the swing points are a good indicator for this. We determined a trend by using looking at swingpoints. A swing point is a local extrema of a certain period.

12.2 Where did I start my research

```
#compute swing points
swing_points = swing_points(df['price'], period=n)
```

Parameters of the swing_points function:

- n is the lookback period
- $df['price']$ is the price column of the dataframe (Which is a timeseries)
- P_t is a value of $df['price']$ at time t

We basically got through the time series dataset and look back n P_t values. The highest and lowest points inside of that specific lookback period are the swing points. After that we determine if price is making higher highs or lower lows. We check this by storing the last swing points and wait till we are either making a higher high or a lower low.

12.3 Finding an edge




Figure 2: Visualization of Δ_t outliers and swing points



Figure 3: Visualization of Δ_t outliers and swing points

After plotting different kind of parts of Timeseries Data sets I created I was pretty sure that somewhere I would be finding an edge since the outliers often where at good entry points for a strategy. Only problem was I wasn't able to get any statistical proof of this. After finding a Github Repo which used Vectorbt to test 1000 strategies at the same time

I tried a similar approach and tested different conditions to see if there was some kind of edge.

12.4 Monte Carlo Testing

First of I have to explain how I search for a strategy and how it nearly killed my computer. We start by defining a different kind of sample sizes of all the Δ_t outliers time periods after the outlier detection. We want to know the return after our condition meets and then check in after z minutes after that

```
sampling_percentages = [0.1, 0.2, 0.6, 0.8, 1.0]
holding_periods = [60, 120, 240, 360]
```

The sampling percentage defines how many of the existing outliers we use and test the returns on. So let's say we have x amount of Δ_t Outliers. We then randomly select z % amount of outliers and test the returns on them. We do this 1000 times and then take the average return of all the 1000 tests. The second step is to test each sampling percentage on the different holding periods. So let's say we have a sampling percentage of z % and a holding period of y minutes. We then take the z % amount of outliers and test the returns on them after y minutes. This test is basically copied from the Vectorbt Github Repository. I just changed the code to fit my needs and added some extra information returns. My code gives back the average return, the Sharpe, expectancy and mean Z-Score.

- **Sharpe Ratio:** The Sharpe ratio is a measure of the risk-adjusted return of a strategy. It is calculated by dividing the average return of the strategy by the standard deviation of the returns.
- Calculated by $\frac{\bar{r}}{\sigma}$ Where \bar{r} is the average return and σ is the standard deviation of the returns.
- **Expectancy:** The expectancy is a measure of the profitability of a strategy. It is calculated by dividing the average return of the strategy by the average loss of the strategy.
- Calculated by $\frac{\bar{r}}{\bar{l}}$ Where \bar{r} is the average return and \bar{l} is the average loss.
- **Mean Z-Score:** The mean Z-Score is a measure of the average Z-Score of the strategy. It is calculated by taking the average of the Z-Scores of the strategy.
- Calculated by $\frac{1}{n} \sum_{i=1}^n z_i$ Where z_i is the Z-Score of the strategy at time i and n is the number of Z-Scores.
- **Average Return:** The average return is a measure of the average return of the strategy. It is calculated by taking the average of the returns of the strategy.

- Calculated by $\frac{1}{n} \sum_{i=1}^n r_i$ Where r_i is the return of the strategy at time i and n is the number of returns.

Now it is time to test some strategies. The first strategy I want to test is pretty intuitive. Measure the average return of bullish outliers inside of an uptrend.

spacing

label A outlier is bullish if the the orderbook Δ_t is two standard deviations above the mean of the last 1440 Δ_t values and bearish if it is two standard deviations below the mean of the last 1440 Δ_t values.

Now when we backtest a strategy we have to have a few things in mind. First of all we are backtesting on on historical data and if we just use different kind of entry conditions we might just change the entry conditions till we have a good end results. This would be overfitting and not work on future data. In order to prevent overfitting we have to use a holdout sample. I have one dataset on which we test our entry conditions to see if they even have potential.

If these conditions are met for a condition I test on some out of sample data. (Out of sample data just means that we test on new data) Most of the time things end up not even passing the first out of sample test and lose their potential instantly. If not I have another out of sample test and if that one is passed we are ready to do some more monte carlo testing and look how potential equity curves look like. And to make things even worse we can assume that the sharpen ratio will decrease by atleast 20% compared to the backtest simulation.

12.5 Backtesting framework

My backtesting framework is built in Python using the VectorBT library and consists of several key components:

Data Pipeline

Connection to Postgres database with the following data:

- Historical price data from Coinbase (BTC/USD)
- Orderbook delta data at different depths ($\Delta_{1\%}$, $\Delta_{2.5\%}$, $\Delta_{5\%}$)
- 1-minute timeframe for base calculations
- Direct ema calculation which a lenght of 50 on the 1h timeframe

I choose to directly calculate the EMA on the 1h timeframe since my datamining process might miss random minutes of data so if we calculate the EMA afterwards this could influence the results.

Testing Methodology

The framework implements a three-stage testing process:

1. Initial Sample Testing

- Test strategy on initial dataset
- Multiple sampling percentages: [0.1, 0.2, 0.6, 0.8, 1.0]
- Various holding periods: [60, 120, 240, 360] minutes

2. Out-of-Sample Validation

- Test promising strategies on separate dataset (Two different datasets)
- Require consistent performance across datasets
- Expect 20% degradation in Sharpe ratio

3. Monte Carlo Simulation

- Random sampling of trade opportunities
- 1000 iterations per test configuration
- Analysis of distribution of outcomes

Performance Metrics

Each strategy is evaluated using:

- **Sharpe Ratio:** $\frac{\bar{r}}{\sigma}$
- **Expectancy:** $\frac{\bar{r}}{\bar{l}}$
- **Mean Z-Score:** $\frac{1}{n} \sum_{i=1}^n z_i$
- **Average Return:** $\frac{1}{n} \sum_{i=1}^n r_i$

Risk Management

The framework incorporates:

- Trading fees (0.003% on Hyperliquid)
- Slippage of (0.0001% on Hyperliquid)
- Position sizing rules
- Maximum drawdown limits
- Stop-loss and take-profit levels

Implementation

```
# Example of strategy implementation
def backtest_strategy(data, params):
    # Define entry/exit signals
    outlier_mask = (data['outlier_context'] == 's') &
                  (data['trend'] == 'Uptrend')

    # Create portfolio simulation
    pf = vbt.Portfolio.from_signals(
        close=data['price'],
        entries=outlier_mask,
        exits=exits_after_n_bars(outlier_mask, n=params['holding_period']),
        init_cash=100,
        freq='1T'
    )

    return calculate_metrics(pf)
```

This framework allows for rapid testing of multiple strategy variations while maintaining strict validation criteria to prevent overfitting.

13 Developing a strategy

13.1 Initial thoughts

I will only able to show a few strategies backtests since I tested on about 150 different logics and only 1 of them passed all the out of sample tests. I will show the best ideas I had and the results of the backtests.

13.2 Reverse strategy

So logical thinking we could assume that long strategy can be reversed into a short strategy and vice versa. The problem is that we have fees and slippage. So we need a certain amount of profit to cover the fees and slippage. And since my strategy has a rather high frequency we need to have a lot of trades to cover the fees and slippage.

13.3 First strategy test

So the first strategy I test was before I developed the backtesting framework. I was testing on a simple logic. If the price on the 1h timeframe closed above the 50 period EMA and the Δ_5 of the orderbook was positive we would enter a buy position if price was in a range of 0.05% of the EMA_{50} .

In addition to that I was using a trailing stop. The trailing stop was set to 1% of the price. A trailing stop is a stop loss that is adjusted to the price of the asset. So if the price goes up the stop loss goes up with it. The idea behing this is that we can catch a bigger upwards in contrast to a fixed stop profit.

I only tested this on a short sample of data because I just had started my data mining process and didn't have any more data. But I can assure that more data didn't make it more profitable just worse.

13.3.1 Code

```
# Define entry/exit signals
range_pct = 0.0005
trailing_stop_pct = 0.01

entries = (
    (df['bias'] > df['ema']) &
    (df['price'] >= df['ema'] * (1 - range_pct)) &
    (df['price'] <= df['ema'] * (1 + range_pct))
)
```

13.3.2 Strategy 1 results

Table 4: Detailed backtest results for the first strategy test.

Metric	Value
General	
Start	2025-03-12 00:00:01
End	2025-03-17 07:24:01
Period (minutes)	7,636
Performance	
Start Value	100.00
End Value	97.13
Total Return [%]	-2.87
Benchmark Return [%]	0.38
Total Fees Paid	2.21
Open Trade PnL	0.23
Trades	
Total Trades	23
Closed Trades	22
Win Rate [%]	31.82
Profit Factor	0.75
Expectancy	-0.14
Best Trade [%]	3.15
Worst Trade [%]	-1.24
Avg Winning Trade [%]	1.35
Avg Losing Trade [%]	-0.83

13.4 Explaining figure

So on the Y axis we have the average return of the strategy. And on the X axis we have the sampling percentage. A sampling percentage of 10% means that we are using 10 random % of the outliers and then testing the returns on them. We conduct this 1000 times and then take the average return of all the 1000 tests. In addition to that we are testing the returns on different holding periods. So for example if we are testing the returns on a holding period of 60 minutes we are testing the returns on the 10 random % of the outliers after 60 minutes. We conduct this 1000 times and then take the average return of all the 1000 tests.




Figure 4: Average returns of bullish outliers inside of an uptrend

13.5 What did change

So from march till july I didn't backtest any strategies. I was focusing on developing signals and indicators.

14 sources

14.1 How good or random is your trading

14.2 Testing 1000 strategies at the same time (via Monte Carlo)

14.3 Tweets

A lot of my ideas come from thiy guys Tweets, to research into this. He is a disgresionary trade what means he trades based on his own decisions and has a mental framework to trade. I learned what an orderbook delta is from him.

14.4 TRDR This platform allow you to use different kind of metrics on different Timeseries datasets (BTC/USD Price, orderbookdelta and Open Interest

his platform allow you to use different kind of metrics on different Timeseries datasets like: BTC/USD Price orderbookdelta and Open Interest

14.5 Trend line automation

Used this for the *v3* version of my trend identification system

14.6 Vectorbt github documentation respository.

I implemented some ideas from this repository for my backtesting framework and used vectorbt for the fees and slippage implementations

14.7 Backtesting

I will backtest my strategy on a 1 year dataset from 2024-03-12 to 2025-03-12.